

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: df=pd.read_csv("madrid_2002.csv")
```

```
In [3]: df.head()
```

```
Out[3]:
```

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	PM10	PXY
0	2002-04-01 01:00:00	NaN	1.39	NaN	NaN	NaN	145.100006	352.100006	NaN	6.54	41.990002	NaN
1	2002-04-01 01:00:00	1.93	0.71	2.33	6.2	0.15	98.150002	153.399994	2.67	6.85	20.980000	2.5
2	2002-04-01 01:00:00	NaN	0.80	NaN	NaN	NaN	103.699997	134.000000	NaN	13.01	28.440001	NaN
3	2002-04-01 01:00:00	NaN	1.61	NaN	NaN	NaN	97.599998	268.000000	NaN	5.12	42.180000	NaN
4	2002-04-01 01:00:00	NaN	1.90	NaN	NaN	NaN	92.089996	237.199997	NaN	7.28	76.330002	NaN

```
In [4]: df=df.dropna()
```

```
In [5]: df.columns
```

```
Out[5]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
               'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
              dtype='object')
```

In [6]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 32381 entries, 1 to 217295
Data columns (total 16 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   date        32381 non-null  object  
 1   BEN         32381 non-null  float64  
 2   CO          32381 non-null  float64  
 3   EBE         32381 non-null  float64  
 4   MXY         32381 non-null  float64  
 5   NMHC        32381 non-null  float64  
 6   NO_2        32381 non-null  float64  
 7   NOx         32381 non-null  float64  
 8   OXY         32381 non-null  float64  
 9   O_3         32381 non-null  float64  
10  PM10        32381 non-null  float64  
11  PXY         32381 non-null  float64  
12  SO_2        32381 non-null  float64  
13  TCH         32381 non-null  float64  
14  TOL         32381 non-null  float64  
15  station     32381 non-null  int64  
dtypes: float64(14), int64(1), object(1)
memory usage: 4.2+ MB
```

In [7]: data=df[['CO', 'station']]
data


Out[7]:

	CO	station
1	0.71	28079035
5	0.72	28079006
22	0.80	28079024
24	1.04	28079099
26	0.53	28079035
...
217269	0.28	28079024
217271	1.30	28079099
217273	0.97	28079035
217293	0.58	28079024
217295	1.17	28079099

32381 rows × 2 columns


```
In [8]: data.plot.line(subplots=True)
```

```
Out[8]: array([<Axes: >, <Axes: >], dtype=object)
```




```
In [9]: b=data[0:50]  
b.plot.bar()
```

Out[9]: <Axes: >




```
In [10]: data.plot.hist()
```

```
Out[10]: <Axes: ylabel='Frequency'>
```




```
In [11]: data.plot.area()
```

```
Out[11]: <Axes: >
```




```
In [12]: data.plot.box()
```

```
Out[12]: <Axes: >
```



```
In [13]: data.plot.scatter(x='CO',y='station')
```

```
Out[13]: <Axes: xlabel='CO', ylabel='station'>
```




```
In [14]: x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
             'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]  
y=df['station']
```

```
In [15]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

Linear Regression


```
In [16]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
lr.intercept_
prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[16]: <matplotlib.collections.PathCollection at 0x1491a728910>



```
In [17]: print(lr.score(x_test,y_test))
print(lr.score(x_train,y_train))
```

0.19808264306622458

0.1988584434706827

Ridge and Lasso

```
In [18]: from sklearn.linear_model import Ridge,Lasso
rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
print(rr.score(x_test,y_test))
print(rr.score(x_train,y_train))
la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

0.19746964338217032

0.1986451008739084

Out[18]: Lasso(alpha=10)

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [19]: la.score(x_test,y_test)
```

Out[19]: 0.05959650402235683

ElasticNet

```
In [20]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

Out[20]: ElasticNet()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [21]: en.coef_
```

Out[21]: array([0.92575093, 0. , -2.95775812, 1.56224474, 0.19104727,
 0.22680299, -0.03020663, -2.425336 , -0.02546159, 0.00742185,
 2.35459024, 0.39870298, 1.05678995, -1.1689577])

```
In [22]: en.intercept_
```

Out[22]: 28079038.632254932

```
In [23]: prediction=en.predict(x_test)
```

```
In [24]: en.score(x_test,y_test)
```

```
Out[24]: 0.1000461075926431
```

Evaluation Metrics

```
In [25]: from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
28.59476137351143
1126.0759677707854
33.55705540971653
```

Logistics Regression

```
In [26]: from sklearn.linear_model import LogisticRegression
```

```
In [27]: feature_matrix=df[['BEN', 'CO', 'EBE', 'MXV', 'NMHC', 'NO_2', 'NOx', 'OXY', 'C
'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
target_vector=df[ 'station']
```

```
In [28]: from sklearn.preprocessing import StandardScaler
fs=StandardScaler().fit_transform(feature_matrix)
logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)
logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)
logr.score(fs,target_vector)
```

```
Out[28]: 0.8480899292795158
```

```
In [29]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
logr.predict_proba(observation)
```

```
Out[29]: array([[2.53067744e-10, 3.44181204e-71, 1.00000000e+00, 1.42366623e-13]])
```

Random Forest

```
In [30]: from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

Out[30]: RandomForestClassifier()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [31]: parameters={'max_depth':[1,2,3,4,5],
'min_samples_leaf':[5,10,15,20,25],
'n_estimators':[10,20,30,40,50]
}
```

```
In [32]: from sklearn.model_selection import GridSearchCV
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

Out[32]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
param_grid={'max_depth': [1, 2, 3, 4, 5],
'min_samples_leaf': [5, 10, 15, 20, 25],
'n_estimators': [10, 20, 30, 40, 50]},
scoring='accuracy')

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [33]: rfc_best=grid_search.best_estimator_  
         from sklearn.tree import plot_tree  
         plt.figure(figsize=(80,40))  
         plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b'])
```

```

Out[33]: [Text(0.4639830508474576, 0.9166666666666666, 'PXY <= 1.005\ngini = 0.749\nsa
mples = 14380\nvalue = [5122, 5733, 5773, 6038]\nnclass = d'),
Text(0.19915254237288135, 0.75, 'MXY <= 1.105\ngini = 0.413\nsamples = 3283
\nvalue = [106, 3916, 736, 484]\nnclass = b'),
Text(0.07627118644067797, 0.5833333333333334, 'TCH <= 1.175\ngini = 0.151\nsa
mples = 1605\nvalue = [9, 2354, 158, 40]\nnclass = b'),
Text(0.03389830508474576, 0.4166666666666667, 'SO_2 <= 6.66\ngini = 0.305\nsa
mples = 67\nvalue = [8, 11, 89, 0]\nnclass = c'),
Text(0.01694915254237288, 0.25, 'gini = 0.658\nsamples = 15\nvalue = [8, 11,
8, 0]\nnclass = b'),
Text(0.05084745762711865, 0.25, 'gini = 0.0\nsamples = 52\nvalue = [0, 0, 8
1, 0]\nnclass = c'),
Text(0.11864406779661017, 0.4166666666666667, 'NOx <= 89.385\ngini = 0.087\n
samples = 1538\nvalue = [1, 2343, 69, 40]\nnclass = b'),
Text(0.0847457627118644, 0.25, 'TCH <= 1.225\ngini = 0.068\nsamples = 1494\n
value = [1, 2286, 45, 37]\nnclass = b'),
Text(0.06779661016949153, 0.08333333333333333, 'gini = 0.442\nsamples = 106
\nvalue = [1, 122, 36, 12]\nnclass = b'),
Text(0.1016949152542373, 0.08333333333333333, 'gini = 0.031\nsamples = 1388
\nvalue = [0, 2164, 9, 25]\nnclass = b'),
Text(0.15254237288135594, 0.25, 'O_3 <= 7.23\ngini = 0.457\nsamples = 44\nva
lue = [0, 57, 24, 3]\nnclass = b'),
Text(0.13559322033898305, 0.08333333333333333, 'gini = 0.074\nsamples = 27\n
value = [0, 50, 2, 0]\nnclass = b'),
Text(0.1694915254237288, 0.08333333333333333, 'gini = 0.471\nsamples = 17\nv
alue = [0, 7, 22, 3]\nnclass = c'),
Text(0.3220338983050847, 0.5833333333333334, 'NMHC <= 0.065\ngini = 0.585\nsa
mples = 1678\nvalue = [97, 1562, 578, 444]\nnclass = b'),
Text(0.2542372881355932, 0.4166666666666667, 'TOL <= 2.435\ngini = 0.55\nsam
ples = 516\nvalue = [91, 30, 503, 191]\nnclass = c'),
Text(0.22033898305084745, 0.25, 'NMHC <= 0.035\ngini = 0.477\nsamples = 101
\nvalue = [3, 4, 45, 101]\nnclass = d'),
Text(0.2033898305084746, 0.08333333333333333, 'gini = 0.475\nsamples = 37\nv
alue = [3, 0, 36, 15]\nnclass = c'),
Text(0.23728813559322035, 0.08333333333333333, 'gini = 0.235\nsamples = 64\n
value = [0, 4, 9, 86]\nnclass = d'),
Text(0.288135593220339, 0.25, 'NO_2 <= 19.82\ngini = 0.484\nsamples = 415\nv
alue = [88, 26, 458, 90]\nnclass = c'),
Text(0.2711864406779661, 0.08333333333333333, 'gini = 0.715\nsamples = 56\nv
alue = [26, 21, 33, 10]\nnclass = c'),
Text(0.3050847457627119, 0.08333333333333333, 'gini = 0.417\nsamples = 359\n
value = [62, 5, 425, 80]\nnclass = c'),
Text(0.3898305084745763, 0.4166666666666667, 'NOx <= 80.985\ngini = 0.306\nsa
mples = 1162\nvalue = [6, 1532, 75, 253]\nnclass = b'),
Text(0.3559322033898305, 0.25, 'PXY <= 0.785\ngini = 0.24\nsamples = 990\nva
lue = [1, 1391, 31, 187]\nnclass = b'),
Text(0.3389830508474576, 0.08333333333333333, 'gini = 0.101\nsamples = 681\n
value = [0, 1046, 15, 43]\nnclass = b'),
Text(0.3728813559322034, 0.08333333333333333, 'gini = 0.453\nsamples = 309\n
value = [1, 345, 16, 144]\nnclass = b'),
Text(0.423728813559322, 0.25, 'NO_2 <= 62.125\ngini = 0.6\nsamples = 172\nva
lue = [5, 141, 44, 66]\nnclass = b'),
Text(0.4067796610169492, 0.08333333333333333, 'gini = 0.53\nsamples = 49\nva
lue = [1, 18, 8, 46]\nnclass = d'),
Text(0.4406779661016949, 0.08333333333333333, 'gini = 0.497\nsamples = 123\n
value = [4, 123, 36, 20]\nnclass = b'),
Text(0.7288135593220338, 0.75, 'OXY <= 4.785\ngini = 0.721\nsamples = 11097

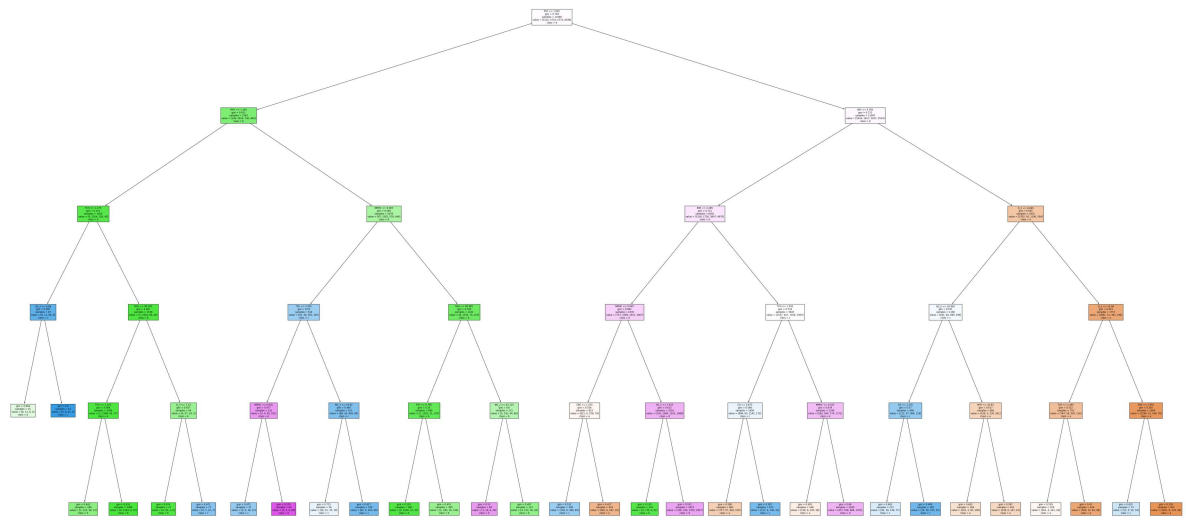
```

```

\value = [5016, 1817, 5037, 5554]\nclass = d'),
Text(0.5932203389830508, 0.5833333333333334, 'EBE <= 2.285\ngini = 0.711\nsa
mples = 8182\nvalue = [2314, 1726, 3847, 4970]\nclass = d'),
Text(0.5254237288135594, 0.4166666666666667, 'NMHC <= 0.045\ngini = 0.682\ns
amples = 4339\nvalue = [757, 1099, 1931, 3067]\nclass = d'),
Text(0.4915254237288136, 0.25, 'EBE <= 1.555\ngini = 0.556\nsamples = 823\nv
alue = [621, 0, 576, 79]\nclass = a'),
Text(0.4745762711864407, 0.0833333333333333, 'gini = 0.521\nsamples = 399\n
value = [156, 0, 389, 69]\nclass = c'),
Text(0.5084745762711864, 0.0833333333333333, 'gini = 0.427\nsamples = 424\n
value = [465, 0, 187, 10]\nclass = a'),
Text(0.559322033898305, 0.25, 'SO_2 <= 5.625\ngini = 0.615\nsamples = 3516\n
value = [136, 1099, 1355, 2988]\nclass = d'),
Text(0.5423728813559322, 0.0833333333333333, 'gini = 0.135\nsamples = 543\n
value = [0, 779, 0, 61]\nclass = b'),
Text(0.576271186440678, 0.0833333333333333, 'gini = 0.531\nsamples = 2973\n
value = [136, 320, 1355, 2927]\nclass = d'),
Text(0.6610169491525424, 0.4166666666666667, 'TCH <= 1.345\ngini = 0.719\nsa
mples = 3843\nvalue = [1557, 627, 1916, 1903]\nclass = c'),
Text(0.6271186440677966, 0.25, 'CO <= 0.675\ngini = 0.586\nsamples = 1495\nv
alue = [994, 63, 1142, 172]\nclass = c'),
Text(0.6101694915254238, 0.0833333333333333, 'gini = 0.588\nsamples = 862\n
value = [777, 57, 402, 150]\nclass = a'),
Text(0.6440677966101694, 0.0833333333333333, 'gini = 0.387\nsamples = 633\n
value = [217, 6, 740, 22]\nclass = c'),
Text(0.6949152542372882, 0.25, 'NMHC <= 0.105\ngini = 0.679\nsamples = 2348
\nvalue = [563, 564, 774, 1731]\nclass = d'),
Text(0.6779661016949152, 0.0833333333333333, 'gini = 0.605\nsamples = 168\n
value = [126, 6, 105, 28]\nclass = a'),
Text(0.711864406779661, 0.0833333333333333, 'gini = 0.66\nsamples = 2180\nv
alue = [437, 558, 669, 1703]\nclass = d'),
Text(0.864406779661017, 0.5833333333333334, 'O_3 <= 8.695\ngini = 0.565\nsam
ples = 2915\nvalue = [2702, 91, 1190, 584]\nclass = a'),
Text(0.7966101694915254, 0.4166666666666667, 'SO_2 <= 20.965\ngini = 0.671\n
samples = 1158\nvalue = [641, 60, 699, 390]\nclass = c'),
Text(0.7627118644067796, 0.25, 'CO <= 1.335\ngini = 0.57\nsamples = 498\nval
ue = [111, 57, 469, 129]\nclass = c'),
Text(0.7457627118644068, 0.0833333333333333, 'gini = 0.683\nsamples = 213\n
value = [92, 19, 136, 77]\nclass = c'),
Text(0.7796610169491526, 0.0833333333333333, 'gini = 0.409\nsamples = 285\n
value = [19, 38, 333, 52]\nclass = c'),
Text(0.8305084745762712, 0.25, 'MXY <= 20.83\ngini = 0.617\nsamples = 660\nv
alue = [530, 3, 230, 261]\nclass = a'),
Text(0.8135593220338984, 0.0833333333333333, 'gini = 0.605\nsamples = 369\n
value = [291, 3, 83, 200]\nclass = a'),
Text(0.847457627118644, 0.0833333333333333, 'gini = 0.587\nsamples = 291\nv
alue = [239, 0, 147, 61]\nclass = a'),
Text(0.9322033898305084, 0.4166666666666667, 'O_3 <= 18.08\ngini = 0.413\nsa
mples = 1757\nvalue = [2061, 31, 491, 194]\nclass = a'),
Text(0.8983050847457628, 0.25, 'TCH <= 1.465\ngini = 0.522\nsamples = 753\nv
alue = [767, 18, 305, 116]\nclass = a'),
Text(0.8813559322033898, 0.0833333333333333, 'gini = 0.539\nsamples = 329\n
value = [261, 4, 241, 18]\nclass = a'),
Text(0.9152542372881356, 0.0833333333333333, 'gini = 0.42\nsamples = 424\nv
alue = [506, 14, 64, 98]\nclass = a'),
Text(0.9661016949152542, 0.25, 'EBE <= 3.855\ngini = 0.305\nsamples = 1004\n
value = [1294, 13, 186, 78]\nclass = a'),

```

```
Text(0.9491525423728814, 0.08333333333333333, 'gini = 0.625\nsamples = 73\nvalue = [13, 5, 52, 34]\nnclass = c'),
Text(0.9830508474576272, 0.08333333333333333, 'gini = 0.228\nsamples = 931\nvalue = [1281, 8, 134, 44]\nnclass = a')]
```



Conclusion

```
In [34]: print("Linear Regression:",lr.score(x_test,y_test))
print("Ridge Regression:",rr.score(x_test,y_test))
print("Lasso Regression",la.score(x_test,y_test))
print("ElasticNet Regression:",en.score(x_test,y_test))
print("Logistic Regression:",logr.score(fs,target_vector))
print("Random Forest:",grid_search.best_score_)
```

```
Linear Regression: 0.19808264306622458
Ridge Regression: 0.19746964338217032
Lasso Regression 0.05959650402235683
ElasticNet Regression: 0.1000461075926431
Logistic Regression: 0.8480899292795158
Random Forest: 0.7738462895967528
```

Logistic Is Better!!!