

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: df=pd.read_csv("madrid_2008.csv")
```

```
In [3]: df.head()
```

Out[3]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PM10
0	2008-06-01 01:00:00	NaN	0.47	NaN	NaN	NaN	83.089996	120.699997	NaN	16.990000	16.889999
1	2008-06-01 01:00:00	NaN	0.59	NaN	NaN	NaN	94.820000	130.399994	NaN	17.469999	19.040001
2	2008-06-01 01:00:00	NaN	0.55	NaN	NaN	NaN	75.919998	104.599998	NaN	13.470000	20.270000
3	2008-06-01 01:00:00	NaN	0.36	NaN	NaN	NaN	61.029999	66.559998	NaN	23.110001	10.850000
4	2008-06-01 01:00:00	1.68	0.80	1.7	3.01	0.3	105.199997	214.899994	1.61	12.120000	37.160000

```
In [4]: df=df.dropna()
```

```
In [5]: df.columns
```

Out[5]: Index(['date', 'BEN', 'CO', 'EBE', 'MXV', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
dtype='object')

In [6]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 25631 entries, 4 to 226391
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0    date      25631 non-null  object
1    BEN       25631 non-null  float64
2    CO        25631 non-null  float64
3    EBE       25631 non-null  float64
4    MXY       25631 non-null  float64
5    NMHC      25631 non-null  float64
6    NO_2      25631 non-null  float64
7    NOx       25631 non-null  float64
8    OXY       25631 non-null  float64
9    O_3       25631 non-null  float64
10   PM10      25631 non-null  float64
11   PM25      25631 non-null  float64
12   PXY       25631 non-null  float64
13   SO_2      25631 non-null  float64
14   TCH       25631 non-null  float64
15   TOL       25631 non-null  float64
16   station   25631 non-null  int64
dtypes: float64(15), int64(1), object(1)
memory usage: 3.5+ MB
```

In [7]: data=df[['CO', 'station']]
data


Out[7]:

	CO	station
4	0.80	28079006
21	0.37	28079024
25	0.39	28079099
30	0.51	28079006
47	0.39	28079024
...
226362	0.35	28079024
226366	0.46	28079099
226371	0.53	28079006
226387	0.30	28079024
226391	0.36	28079099

25631 rows × 2 columns


```
In [8]: data.plot.line(subplots=True)
```

```
Out[8]: array([<Axes: >, <Axes: >], dtype=object)
```




```
In [9]: b=data[0:50]  
b.plot.bar()
```

Out[9]: <Axes: >




```
In [10]: data.plot.hist()
```

```
Out[10]: <Axes: ylabel='Frequency'>
```




```
In [11]: data.plot.area()
```

```
Out[11]: <Axes: >
```




```
In [12]: data.plot.box()
```

```
Out[12]: <Axes: >
```



```
In [13]: data.plot.scatter(x='CO',y='station')
```

```
Out[13]: <Axes: xlabel='CO', ylabel='station'>
```




```
In [14]: x=df[['BEN', 'CO', 'EBE', 'MXV', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
              'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]  
y=df['station']
```

```
In [15]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

Linear Regression


```
In [16]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
lr.intercept_
prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[16]: <matplotlib.collections.PathCollection at 0x22bb912c390>



```
In [17]: print(lr.score(x_test,y_test))
print(lr.score(x_train,y_train))
```

0.14189755406565563

0.1442773097866754

Ridge and Lasso

```
In [18]: from sklearn.linear_model import Ridge,Lasso
rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
print(rr.score(x_test,y_test))
print(rr.score(x_train,y_train))
la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

0.14199298931038495

0.14424986480453006

Out[18]:

▼ Lasso

Lasso(alpha=10)

```
In [19]: la.score(x_test,y_test)
```

Out[19]: 0.042320117094826304

ElasticNet

```
In [20]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

Out[20]:

▼ ElasticNet

ElasticNet()

```
In [21]: en.coef_
```

Out[21]: array([-4.75245423, -0. , 0. , 3.34582701, -0. ,
 0.04525883, 0.03148858, 1.44837803, -0.14826515, 0.13088534,
 1.59485819, -0.91313965, 0. , -2.56737519])

```
In [22]: en.intercept_
```

Out[22]: 28079056.835846838

```
In [23]: prediction=en.predict(x_test)
```

```
In [24]: en.score(x_test,y_test)
```

Out[24]: 0.09945769268475402

Evaluation Metrics

```
In [25]: from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
35.734393493612195
1488.6577699930538
38.58312804831995
```

Logistics Regression

```
In [26]: from sklearn.linear_model import LogisticRegression
```

```
In [27]: feature_matrix=df[['BEN', 'CO', 'EBE', 'MXV', 'NMHC', 'NO_2', 'NOx', 'OXY', 'C
'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
target_vector=df[ 'station']
```

```
In [28]: from sklearn.preprocessing import StandardScaler
fs=StandardScaler().fit_transform(feature_matrix)
logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)
logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)
logr.score(fs,target_vector)
```

```
Out[28]: 0.794194530061254
```

```
In [29]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
logr.predict_proba(observation)
```

```
Out[29]: array([[8.32180727e-09, 1.19114483e-13, 9.99999992e-01]])
```

Random Forest

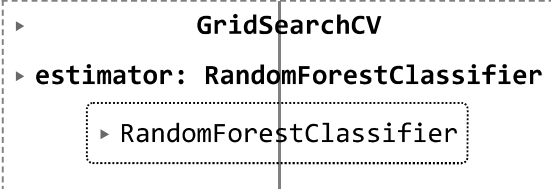
```
In [30]: from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[30]: ▾ RandomForestClassifier
RandomForestClassifier()
```

```
In [31]: parameters={'max_depth':[1,2,3,4,5],  
                  'min_samples_leaf':[5,10,15,20,25],  
                  'n_estimators':[10,20,30,40,50]  
                }
```

```
In [32]: from sklearn.model_selection import GridSearchCV  
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="ac  
grid_search.fit(x_train,y_train)
```

```
Out[32]:
```



```
▶ GridSearchCV  
▶ estimator: RandomForestClassifier  
    ▶ RandomForestClassifier
```

```
In [33]: rfc_best=grid_search.best_estimator_  
         from sklearn.tree import plot_tree  
         plt.figure(figsize=(80,40))  
         plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b'])
```

```

Out[33]: [Text(0.5, 0.9166666666666666, 'NO_2 <= 23.235\ngini = 0.667\nsamples = 11358\nvalue = [6112, 5866, 5963]\nclass = a'),
Text(0.25, 0.75, 'TOL <= 1.005\ngini = 0.276\nsamples = 2561\nvalue = [206, 3440, 438]\nclass = b'),
Text(0.125, 0.5833333333333334, 'PXY <= 0.675\ngini = 0.072\nsamples = 1327\nvalue = [18, 2062, 62]\nclass = b'),
Text(0.0625, 0.4166666666666667, 'SO_2 <= 7.755\ngini = 0.338\nsamples = 219\nvalue = [17, 302, 60]\nclass = b'),
Text(0.03125, 0.25, 'PM10 <= 11.235\ngini = 0.55\nsamples = 85\nvalue = [13, 84, 50]\nclass = b'),
Text(0.015625, 0.08333333333333333, 'gini = 0.533\nsamples = 52\nvalue = [5, 33, 50]\nclass = c'),
Text(0.046875, 0.08333333333333333, 'gini = 0.234\nsamples = 33\nvalue = [8, 51, 0]\nclass = b'),
Text(0.09375, 0.25, 'NOx <= 25.055\ngini = 0.115\nsamples = 134\nvalue = [4, 218, 10]\nclass = b'),
Text(0.078125, 0.08333333333333333, 'gini = 0.046\nsamples = 124\nvalue = [1, 210, 4]\nclass = b'),
Text(0.109375, 0.08333333333333333, 'gini = 0.623\nsamples = 10\nvalue = [3, 8, 6]\nclass = b'),
Text(0.1875, 0.4166666666666667, 'NO_2 <= 21.195\ngini = 0.003\nsamples = 1108\nvalue = [1, 1760, 2]\nclass = b'),
Text(0.15625, 0.25, 'EBE <= 0.555\ngini = 0.002\nsamples = 1081\nvalue = [1, 1723, 1]\nclass = b'),
Text(0.140625, 0.08333333333333333, 'gini = 0.022\nsamples = 63\nvalue = [0, 91, 1]\nclass = b'),
Text(0.171875, 0.08333333333333333, 'gini = 0.001\nsamples = 1018\nvalue = [1, 1632, 0]\nclass = b'),
Text(0.21875, 0.25, 'BEN <= 0.32\ngini = 0.051\nsamples = 27\nvalue = [0, 37, 1]\nclass = b'),
Text(0.203125, 0.08333333333333333, 'gini = 0.0\nsamples = 17\nvalue = [0, 22, 0]\nclass = b'),
Text(0.234375, 0.08333333333333333, 'gini = 0.117\nsamples = 10\nvalue = [0, 15, 1]\nclass = b'),
Text(0.375, 0.5833333333333334, 'NMHC <= 0.195\ngini = 0.45\nsamples = 1234\nvalue = [188, 1378, 376]\nclass = b'),
Text(0.3125, 0.4166666666666667, 'BEN <= 0.485\ngini = 0.641\nsamples = 579\nvalue = [180, 356, 357]\nclass = c'),
Text(0.28125, 0.25, 'OXY <= 0.995\ngini = 0.604\nsamples = 459\nvalue = [91, 282, 322]\nclass = c'),
Text(0.265625, 0.08333333333333333, 'gini = 0.486\nsamples = 285\nvalue = [87, 50, 287]\nclass = c'),
Text(0.296875, 0.08333333333333333, 'gini = 0.25\nsamples = 174\nvalue = [4, 232, 35]\nclass = b'),
Text(0.34375, 0.25, 'PXY <= 0.615\ngini = 0.627\nsamples = 120\nvalue = [89, 74, 35]\nclass = a'),
Text(0.328125, 0.08333333333333333, 'gini = 0.299\nsamples = 53\nvalue = [77, 9, 7]\nclass = a'),
Text(0.359375, 0.08333333333333333, 'gini = 0.533\nsamples = 67\nvalue = [12, 65, 28]\nclass = b'),
Text(0.4375, 0.4166666666666667, 'NO_2 <= 20.225\ngini = 0.05\nsamples = 655\nvalue = [8, 1022, 19]\nclass = b'),
Text(0.40625, 0.25, 'OXY <= 0.775\ngini = 0.014\nsamples = 551\nvalue = [1, 873, 5]\nclass = b'),
Text(0.390625, 0.08333333333333333, 'gini = 0.123\nsamples = 47\nvalue = [1, 72, 4]\nclass = b'),
Text(0.421875, 0.08333333333333333, 'gini = 0.002\nsamples = 504\nvalue =

```

```

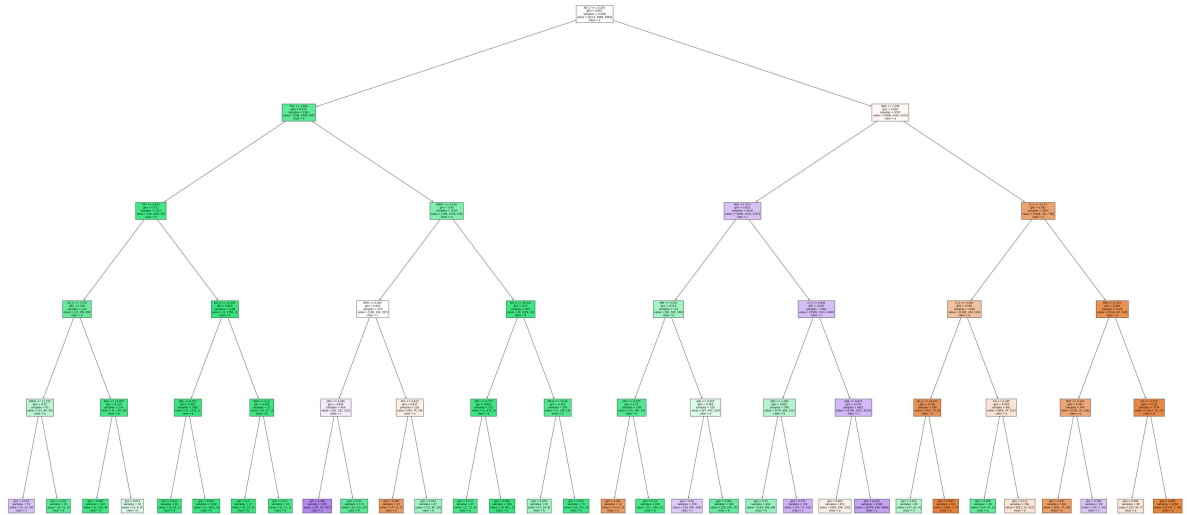
[0, 801, 1]\n\nclass = b'),
  Text(0.46875, 0.25, 'PM10 <= 10.09\ngini = 0.223\nsamples = 104\nvalue = [7,
149, 14]\n\nclass = b'),
  Text(0.453125, 0.08333333333333333, 'gini = 0.559\nsamples = 25\nvalue = [7,
24, 9]\n\nclass = b'),
  Text(0.484375, 0.08333333333333333, 'gini = 0.074\nsamples = 79\nvalue = [0,
125, 5]\n\nclass = b'),
  Text(0.75, 0.75, 'BEN <= 1.385\ngini = 0.629\nsamples = 8797\nvalue = [5906,
2426, 5525]\n\nclass = a'),
  Text(0.625, 0.5833333333333334, 'NOx <= 37.4\ngini = 0.625\nsamples = 6100\n
value = [2600, 2224, 4787]\n\nclass = c'),
  Text(0.5625, 0.4166666666666667, 'EBE <= 0.555\ngini = 0.519\nsamples = 719
\nvalue = [92, 703, 338]\n\nclass = b'),
  Text(0.53125, 0.25, 'OXY <= 0.355\ngini = 0.21\nsamples = 209\nvalue = [25,
300, 14]\n\nclass = b'),
  Text(0.515625, 0.08333333333333333, 'gini = 0.291\nsamples = 11\nvalue = [1
4, 0, 3]\n\nclass = a'),
  Text(0.546875, 0.08333333333333333, 'gini = 0.13\nsamples = 198\nvalue = [1
1, 300, 11]\n\nclass = b'),
  Text(0.59375, 0.25, 'EBE <= 0.955\ngini = 0.569\nsamples = 510\nvalue = [67,
403, 324]\n\nclass = b'),
  Text(0.578125, 0.08333333333333333, 'gini = 0.55\nsamples = 350\nvalue = [3
9, 208, 298]\n\nclass = c'),
  Text(0.609375, 0.08333333333333333, 'gini = 0.363\nsamples = 160\nvalue = [2
8, 195, 26]\n\nclass = b'),
  Text(0.6875, 0.4166666666666667, 'O_3 <= 8.995\ngini = 0.605\nsamples = 5381
\nvalue = [2508, 1521, 4449]\n\nclass = c'),
  Text(0.65625, 0.25, 'OXY <= 1.005\ngini = 0.602\nsamples = 560\nvalue = [17
0, 469, 235]\n\nclass = b'),
  Text(0.640625, 0.08333333333333333, 'gini = 0.53\nsamples = 404\nvalue = [13
6, 396, 94]\n\nclass = b'),
  Text(0.671875, 0.08333333333333333, 'gini = 0.571\nsamples = 156\nvalue = [3
4, 73, 141]\n\nclass = c'),
  Text(0.71875, 0.25, 'EBE <= 0.615\ngini = 0.579\nsamples = 4821\nvalue = [23
38, 1052, 4214]\n\nclass = c'),
  Text(0.703125, 0.08333333333333333, 'gini = 0.603\nsamples = 623\nvalue = [4
59, 396, 128]\n\nclass = a'),
  Text(0.734375, 0.08333333333333333, 'gini = 0.529\nsamples = 4198\nvalue =
[1879, 656, 4086]\n\nclass = c'),
  Text(0.875, 0.5833333333333334, 'O_3 <= 12.73\ngini = 0.361\nsamples = 2697
\nvalue = [3306, 202, 738]\n\nclass = a'),
  Text(0.8125, 0.4166666666666667, 'O_3 <= 5.595\ngini = 0.505\nsamples = 1169
\nvalue = [1165, 140, 520]\n\nclass = a'),
  Text(0.78125, 0.25, 'SO_2 <= 16.675\ngini = 0.236\nsamples = 368\nvalue = [5
02, 70, 8]\n\nclass = a'),
  Text(0.765625, 0.08333333333333333, 'gini = 0.505\nsamples = 65\nvalue = [3
7, 63, 4]\n\nclass = b'),
  Text(0.796875, 0.08333333333333333, 'gini = 0.045\nsamples = 303\nvalue = [4
65, 7, 4]\n\nclass = a'),
  Text(0.84375, 0.25, 'CO <= 0.385\ngini = 0.544\nsamples = 801\nvalue = [663,
70, 512]\n\nclass = a'),
  Text(0.828125, 0.08333333333333333, 'gini = 0.289\nsamples = 38\nvalue = [1
0, 47, 0]\n\nclass = b'),
  Text(0.859375, 0.08333333333333333, 'gini = 0.512\nsamples = 763\nvalue = [6
53, 23, 512]\n\nclass = a'),
  Text(0.9375, 0.4166666666666667, 'BEN <= 1.715\ngini = 0.209\nsamples = 1528
\nvalue = [2141, 62, 218]\n\nclass = a'),

```

```

Text(0.90625, 0.25, 'MXY <= 4.545\ngini = 0.391\nsamples = 454\nvalue = [52
9, 29, 146]\nnclass = a'),
Text(0.890625, 0.08333333333333333, 'gini = 0.327\nsamples = 391\nvalue = [4
91, 27, 92]\nnclass = a'),
Text(0.921875, 0.08333333333333333, 'gini = 0.506\nsamples = 63\nvalue = [3
8, 2, 54]\nnclass = c'),
Text(0.96875, 0.25, 'CO <= 0.375\ngini = 0.116\nsamples = 1074\nvalue = [161
2, 33, 72]\nnclass = a'),
Text(0.953125, 0.08333333333333333, 'gini = 0.584\nsamples = 38\nvalue = [3
3, 26, 7]\nnclass = a'),
Text(0.984375, 0.08333333333333333, 'gini = 0.084\nsamples = 1036\nvalue =
[1579, 7, 65]\nnclass = a'))]

```



Conclusion

```

In [34]: print("Linear Regression:",lr.score(x_test,y_test))
print("Ridge Regression:",rr.score(x_test,y_test))
print("Lasso Regression",la.score(x_test,y_test))
print("ElasticNet Regression:",en.score(x_test,y_test))
print("Logistic Regression:",logr.score(fs,target_vector))
print("Random Forest:",grid_search.best_score_)

```

```

Linear Regression: 0.14189755406565563
Ridge Regression: 0.14199298931038495
Lasso Regression 0.042320117094826304
ElasticNet Regression: 0.09945769268475402
Logistic Regression: 0.794194530061254
Random Forest: 0.852460771963469

```

Logistic Is Better!!!

