

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: df=pd.read_csv("madrid_2011.csv")
```

```
In [3]: df.head()
```

Out[3]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	sta
0	2011-11-01 01:00:00	NaN	1.0	NaN	NaN	154.0	84.0	NaN	NaN	NaN	6.0	NaN	NaN	28079
1	2011-11-01 01:00:00	2.5	0.4	3.5	0.26	68.0	92.0	3.0	40.0	24.0	9.0	1.54	8.7	28079
2	2011-11-01 01:00:00	2.9	NaN	3.8	NaN	96.0	99.0	NaN	NaN	NaN	NaN	NaN	7.2	28079
3	2011-11-01 01:00:00	NaN	0.6	NaN	NaN	60.0	83.0	2.0	NaN	NaN	NaN	NaN	NaN	28079
4	2011-11-01 01:00:00	NaN	NaN	NaN	NaN	44.0	62.0	3.0	NaN	NaN	3.0	NaN	NaN	28079

```
In [4]: df=df.dropna()
```

```
In [5]: df.columns
```

Out[5]: Index(['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO\_2', 'O\_3', 'PM10', 'PM25', 'SO\_2', 'TCH', 'TOL', 'station'], dtype='object')

In [6]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 16460 entries, 1 to 209910
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   date        16460 non-null  object  
 1   BEN         16460 non-null  float64  
 2   CO          16460 non-null  float64  
 3   EBE         16460 non-null  float64  
 4   NMHC        16460 non-null  float64  
 5   NO          16460 non-null  float64  
 6   NO_2        16460 non-null  float64  
 7   O_3         16460 non-null  float64  
 8   PM10        16460 non-null  float64  
 9   PM25        16460 non-null  float64  
10   SO_2        16460 non-null  float64  
11   TCH         16460 non-null  float64  
12   TOL         16460 non-null  float64  
13   station     16460 non-null  int64  
dtypes: float64(12), int64(1), object(1)
memory usage: 1.9+ MB
```

In [7]: data=df[['CO','station']]  
data


Out[7]:

	CO	station
1	0.4	28079008
6	0.3	28079024
25	0.3	28079008
30	0.4	28079024
49	0.2	28079008
...	...	...
209862	0.1	28079024
209881	0.1	28079008
209886	0.1	28079024
209905	0.1	28079008
209910	0.1	28079024

16460 rows × 2 columns


```
In [8]: data.plot.line(subplots=True)
```

```
Out[8]: array([<Axes: >, <Axes: >], dtype=object)
```




```
In [9]: b=data[0:50]  
b.plot.bar()
```

Out[9]: <Axes: >




```
In [10]: data.plot.hist()
```

```
Out[10]: <Axes: ylabel='Frequency'>
```




```
In [11]: data.plot.area()
```

```
Out[11]: <Axes: >
```




```
In [12]: data.plot.box()
```

```
Out[12]: <Axes: >
```



```
In [13]: data.plot.scatter(x='CO',y='station')
```

```
Out[13]: <Axes: xlabel='CO', ylabel='station'>
```



```
In [14]: x=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'NO', 'O_3',  
              'PM10','PM25','SO_2', 'TCH', 'TOL']]  
y=df['station']
```


```
In [15]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

## Linear Regression



```
In [16]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
lr.intercept_
prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[16]: <matplotlib.collections.PathCollection at 0x21c55761190>



```
In [17]: print(lr.score(x_test,y_test))
print(lr.score(x_train,y_train))
```

0.6259709949284429

0.6276479573022269

## Ridge and Lasso

```
In [18]: from sklearn.linear_model import Ridge,Lasso
rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
print(rr.score(x_test,y_test))
print(rr.score(x_train,y_train))
la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

0.5877091721395884

0.594469525578609

Out[18]:

▼ Lasso

Lasso(alpha=10)

```
In [19]: la.score(x_test,y_test)
```

Out[19]: 0.2273508727383743

## ElasticNet

```
In [20]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

Out[20]:

▼ ElasticNet

ElasticNet()

```
In [21]: en.coef_
```

Out[21]: array([ 0.2653043 , 0. , -0. , -0. , -0.14076021,  
 0.05067364, -0.04412062, 0.02581064, 0.10485685, -0.16710751,  
 0. , -0.95067196])

```
In [22]: en.intercept_
```

Out[22]: 28079025.14184462

```
In [23]: prediction=en.predict(x_test)
```

```
In [24]: en.score(x_test,y_test)
```

Out[24]: 0.3356452219334174

## Evaluation Metrics

```
In [25]: from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
5.701842582725372
42.51283990709718
6.520187106755233
```

## Logistics Regression

```
In [26]: from sklearn.linear_model import LogisticRegression
```

```
In [27]: feature_matrix=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'NO', 'O_3',
'PM10','PM25','SO_2', 'TCH', 'TOL']]
target_vector=df[ 'station']
```

```
In [28]: from sklearn.preprocessing import StandardScaler
fs=StandardScaler().fit_transform(feature_matrix)
logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)
```

```
Out[28]: LogisticRegression
LogisticRegression(max_iter=10000)
```

```
In [30]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12]]
logr.predict_proba(observation)
```

```
Out[30]: array([[1.00000000e+00, 1.13637857e-16]])
```

## Random Forest

```
In [31]: from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[31]: RandomForestClassifier
RandomForestClassifier()
```

```
In [32]: parameters={'max_depth':[1,2,3,4,5],
'min_samples_leaf':[5,10,15,20,25],
'n_estimators':[10,20,30,40,50]
}
```

```
In [33]: from sklearn.model_selection import GridSearchCV
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="ac
grid_search.fit(x_train,y_train)
```

```
Out[33]:
GridSearchCV
  estimator: RandomForestClassifier
    RandomForestClassifier
```

```
In [34]: rfc_best=grid_search.best_estimator_
from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b']
\nvalue = [01, 0]\nnclass = a'),
Text(0.5813953488372093, 0.25, 'TOL <= 1.55\ngini = 0.017\nsamples = 769
\nvalue = [1239, 11]\nnclass = a'),
Text(0.5581395348837209, 0.08333333333333333, 'gini = 0.107\nsamples = 12
1\nvalue = [182, 11]\nnclass = a'),
Text(0.6046511627906976, 0.08333333333333333, 'gini = 0.0\nsamples = 648
\nvalue = [1057, 0]\nnclass = a'),
Text(0.813953488372093, 0.5833333333333334, 'NO_2 <= 39.5\ngini = 0.434\n
samples = 3589\nvalue = [3893, 1819]\nnclass = a'),
Text(0.7209302325581395, 0.4166666666666667, 'O_3 <= 23.5\ngini = 0.411\n
samples = 555\nvalue = [258, 634]\nnclass = b'),
Text(0.6744186046511628, 0.25, 'NO <= 15.5\ngini = 0.161\nsamples = 211\n
value = [30, 310]\nnclass = b'),
Text(0.6511627906976745, 0.08333333333333333, 'gini = 0.281\nsamples = 96
\nvalue = [27, 133]\nnclass = b'),
Text(0.6976744186046512, 0.08333333333333333, 'gini = 0.033\nsamples = 11
5\nvalue = [3, 177]\nnclass = b'),
Text(0.7674418604651163, 0.25, 'TOL <= 2.55\ngini = 0.485\nsamples = 344
\nvalue = [228, 324]\nnclass = b'),
Text(0.7441860465116279, 0.08333333333333333, 'gini = 0.32\nsamples = 209
\nvalue = [556, 663]\nnclass = a')
plt.show()
```

## Conclusion

```
In [35]: print("Linear Regression:",lr.score(x_test,y_test))
print("Ridge Regression:",rr.score(x_test,y_test))
print("Lasso Regression",la.score(x_test,y_test))
print("ElasticNet Regression:",en.score(x_test,y_test))
print("Logistic Regression:",logr.score(fs,target_vector))
print("Random Forest:",grid_search.best_score_)
```

Linear Regression: 0.6259709949284429  
Ridge Regression: 0.5877091721395884  
Lasso Regression 0.2273508727383743  
ElasticNet Regression: 0.3356452219334174  
Logistic Regression: 0.9262454434993924  
Random Forest: 0.935688248567957

## Random Forest Is Better!!!