

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: df=pd.read_csv("madrid_2016.csv")
```

```
In [3]: df.head()
```

Out[3]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL	sta
0	2016-11-01 01:00:00	NaN	0.7	NaN	NaN	153.0	77.0	NaN	NaN	NaN	7.0	NaN	NaN	28079
1	2016-11-01 01:00:00	3.1	1.1	2.0	0.53	260.0	144.0	4.0	46.0	24.0	18.0	2.44	14.4	28079
2	2016-11-01 01:00:00	5.9	NaN	7.5	NaN	297.0	139.0	NaN	NaN	NaN	NaN	NaN	26.0	28079
3	2016-11-01 01:00:00	NaN	1.0	NaN	NaN	154.0	113.0	2.0	NaN	NaN	NaN	NaN	NaN	28079
4	2016-11-01 01:00:00	NaN	NaN	NaN	NaN	275.0	127.0	2.0	NaN	NaN	18.0	NaN	NaN	28079

```
In [4]: df=df.dropna()
```

```
In [5]: df.columns
```

Out[5]: Index(['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO\_2', 'O\_3', 'PM10', 'PM25', 'SO\_2', 'TCH', 'TOL', 'station'], dtype='object')

In [6]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 16932 entries, 1 to 209478
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   date        16932 non-null  object  
 1   BEN         16932 non-null  float64  
 2   CO          16932 non-null  float64  
 3   EBE         16932 non-null  float64  
 4   NMHC        16932 non-null  float64  
 5   NO          16932 non-null  float64  
 6   NO_2        16932 non-null  float64  
 7   O_3         16932 non-null  float64  
 8   PM10        16932 non-null  float64  
 9   PM25        16932 non-null  float64  
10   SO_2        16932 non-null  float64  
11   TCH         16932 non-null  float64  
12   TOL         16932 non-null  float64  
13   station     16932 non-null  int64  
dtypes: float64(12), int64(1), object(1)
memory usage: 1.9+ MB
```

In [7]: `data=df[['CO','station']]`  
`data`


Out[7]:

	CO	station
1	1.1	28079008
6	0.8	28079024
25	1.0	28079008
30	0.7	28079024
49	0.8	28079008
...	...	...
209430	0.2	28079024
209449	0.4	28079008
209454	0.2	28079024
209473	0.4	28079008
209478	0.2	28079024

16932 rows × 2 columns


```
In [8]: data.plot.line(subplots=True)
```

```
Out[8]: array([<Axes: >, <Axes: >], dtype=object)
```




```
In [9]: b=data[0:50]  
b.plot.bar()
```

Out[9]: <Axes: >




```
In [10]: data.plot.hist()
```

```
Out[10]: <Axes: ylabel='Frequency'>
```




```
In [11]: data.plot.area()
```

```
Out[11]: <Axes: >
```




```
In [12]: data.plot.box()
```

```
Out[12]: <Axes: >
```



```
In [13]: data.plot.scatter(x='CO',y='station')
```

```
Out[13]: <Axes: xlabel='CO', ylabel='station'>
```



```
In [14]: x=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'NO', 'O_3',  
              'PM10','PM25','SO_2', 'TCH', 'TOL']]  
y=df['station']
```


```
In [15]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

## Linear Regression



```
In [16]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
lr.intercept_
prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[16]: <matplotlib.collections.PathCollection at 0x1f37805afd0>



```
In [17]: print(lr.score(x_test,y_test))
print(lr.score(x_train,y_train))
```

0.8307069955786595  
0.8268326961423023

## Ridge and Lasso

```
In [18]: from sklearn.linear_model import Ridge,Lasso
rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
print(rr.score(x_test,y_test))
print(rr.score(x_train,y_train))
la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

0.8307958700138998

0.8267460978017672

Out[18]:

▼ Lasso

Lasso(alpha=10)

```
In [19]: la.score(x_test,y_test)
```

Out[19]: 0.6552591684765237

## ElasticNet

```
In [20]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

Out[20]:

▼ ElasticNet

ElasticNet()

```
In [21]: en.coef_
```

Out[21]: array([-0. , 0. , -0. , -0. , -0.10673492,  
 0.04673306, -0.02155889, 0.003281 , 0.05059301, -0.86163134,  
 -0.02501396, 0. ])

```
In [22]: en.intercept_
```

Out[22]: 28079026.179314196

```
In [23]: prediction=en.predict(x_test)
```

```
In [24]: en.score(x_test,y_test)
```

Out[24]: 0.7176989302087358

## Evaluation Metrics

```
In [25]: from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
3.29590529195715
18.066361126603837
4.250454225915607
```

## Logistics Regression

```
In [26]: from sklearn.linear_model import LogisticRegression
```

```
In [27]: feature_matrix=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'NO', 'O_3',
'PM10','PM25','SO_2', 'TCH', 'TOL']]
target_vector=df[ 'station']
```

```
In [28]: from sklearn.preprocessing import StandardScaler
fs=StandardScaler().fit_transform(feature_matrix)
logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)
```

```
Out[28]: LogisticRegression
LogisticRegression(max_iter=10000)
```

```
In [29]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12]]
logr.predict_proba(observation)
```

```
Out[29]: array([[1.00000000e+00, 1.86372661e-53]])
```

## Random Forest

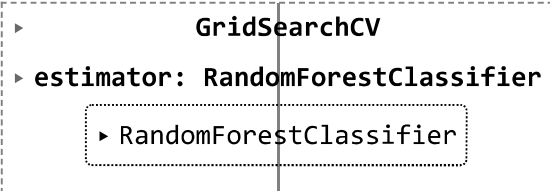
```
In [30]: from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[30]: RandomForestClassifier
RandomForestClassifier()
```

```
In [31]: parameters={'max_depth':[1,2,3,4,5],  
                    'min_samples_leaf':[5,10,15,20,25],  
                    'n_estimators':[10,20,30,40,50]  
                    }
```

```
In [32]: from sklearn.model_selection import GridSearchCV  
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="ac  
grid_search.fit(x_train,y_train)
```

```
Out[32]:
```



```
▶ GridSearchCV  
▶ estimator: RandomForestClassifier  
    ▶ RandomForestClassifier
```

```
In [33]: rfc_best=grid_search.best_estimator_  
         from sklearn.tree import plot_tree  
         plt.figure(figsize=(80,40))  
         plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b'])
```

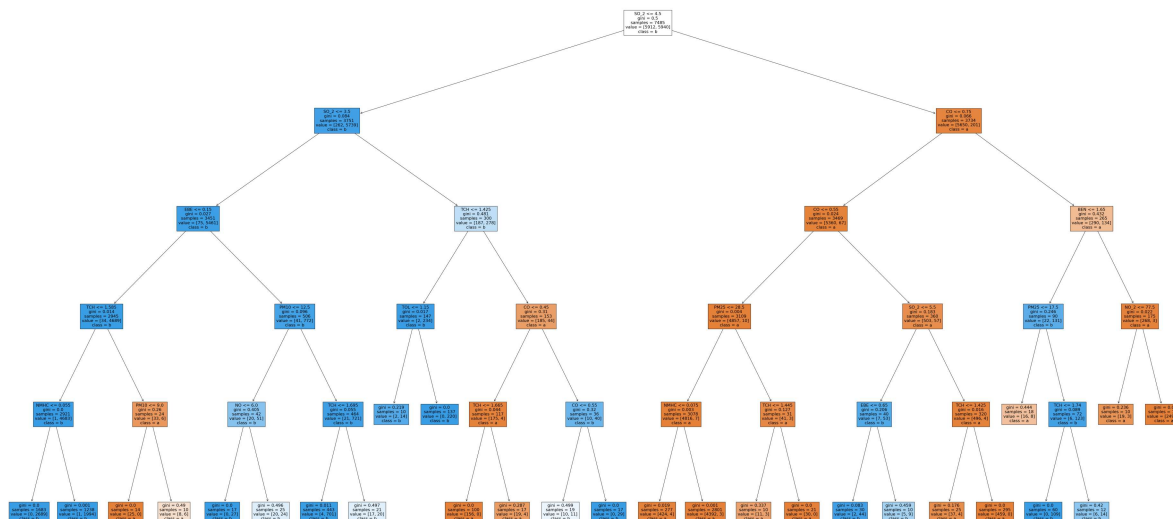
```

Out[33]: [Text(0.5433673469387755, 0.9166666666666666, 'SO_2 <= 4.5\ngini = 0.5\nsamples = 7485\nvalue = [5912, 5940]\nclass = b'),
Text(0.28061224489795916, 0.75, 'SO_2 <= 3.5\ngini = 0.084\nsamples = 3751\nvalue = [262, 5739]\nclass = b'),
Text(0.16326530612244897, 0.5833333333333333, 'EBE <= 0.15\ngini = 0.027\nsamples = 3451\nvalue = [75, 5461]\nclass = b'),
Text(0.08163265306122448, 0.4166666666666667, 'TCH <= 1.585\ngini = 0.014\nsamples = 2945\nvalue = [34, 4689]\nclass = b'),
Text(0.04081632653061224, 0.25, 'NMHC <= 0.055\ngini = 0.0\nsamples = 2921\nvalue = [1, 4683]\nclass = b'),
Text(0.02040816326530612, 0.08333333333333333, 'gini = 0.0\nsamples = 1683\nvalue = [0, 2689]\nclass = b'),
Text(0.061224489795918366, 0.08333333333333333, 'gini = 0.001\nsamples = 1238\nvalue = [1, 1994]\nclass = b'),
Text(0.12244897959183673, 0.25, 'PM10 <= 9.0\ngini = 0.26\nsamples = 24\nvalue = [33, 6]\nclass = a'),
Text(0.10204081632653061, 0.08333333333333333, 'gini = 0.0\nsamples = 14\nvalue = [25, 0]\nclass = a'),
Text(0.14285714285714285, 0.08333333333333333, 'gini = 0.49\nsamples = 10\nvalue = [8, 6]\nclass = a'),
Text(0.24489795918367346, 0.4166666666666667, 'PM10 <= 12.5\ngini = 0.096\nsamples = 506\nvalue = [41, 772]\nclass = b'),
Text(0.20408163265306123, 0.25, 'NO <= 6.0\ngini = 0.405\nsamples = 42\nvalue = [20, 51]\nclass = b'),
Text(0.1836734693877551, 0.08333333333333333, 'gini = 0.0\nsamples = 17\nvalue = [0, 27]\nclass = b'),
Text(0.22448979591836735, 0.08333333333333333, 'gini = 0.496\nsamples = 25\nvalue = [20, 24]\nclass = b'),
Text(0.2857142857142857, 0.25, 'TCH <= 1.695\ngini = 0.055\nsamples = 464\nvalue = [21, 721]\nclass = b'),
Text(0.2653061224489796, 0.08333333333333333, 'gini = 0.011\nsamples = 443\nvalue = [4, 701]\nclass = b'),
Text(0.30612244897959184, 0.08333333333333333, 'gini = 0.497\nsamples = 21\nvalue = [17, 20]\nclass = b'),
Text(0.3979591836734694, 0.5833333333333333, 'TCH <= 1.425\ngini = 0.481\nsamples = 300\nvalue = [187, 278]\nclass = b'),
Text(0.3469387755102041, 0.4166666666666667, 'TOL <= 1.15\ngini = 0.017\nsamples = 147\nvalue = [2, 234]\nclass = b'),
Text(0.32653061224489793, 0.25, 'gini = 0.219\nsamples = 10\nvalue = [2, 14]\nclass = b'),
Text(0.3673469387755102, 0.25, 'gini = 0.0\nsamples = 137\nvalue = [0, 220]\nclass = b'),
Text(0.4489795918367347, 0.4166666666666667, 'CO <= 0.45\ngini = 0.31\nsamples = 153\nvalue = [185, 44]\nclass = a'),
Text(0.40816326530612246, 0.25, 'TCH <= 1.665\ngini = 0.044\nsamples = 117\nvalue = [175, 4]\nclass = a'),
Text(0.3877551020408163, 0.08333333333333333, 'gini = 0.0\nsamples = 100\nvalue = [156, 0]\nclass = a'),
Text(0.42857142857142855, 0.08333333333333333, 'gini = 0.287\nsamples = 17\nvalue = [19, 4]\nclass = a'),
Text(0.4897959183673469, 0.25, 'CO <= 0.55\ngini = 0.32\nsamples = 36\nvalue = [10, 40]\nclass = b'),
Text(0.46938775510204084, 0.08333333333333333, 'gini = 0.499\nsamples = 19\nvalue = [10, 11]\nclass = b'),
Text(0.5102040816326531, 0.08333333333333333, 'gini = 0.0\nsamples = 17\nvalue = [0, 29]\nclass = b'),
Text(0.8061224489795918, 0.75, 'CO <= 0.75\ngini = 0.066\nsamples = 3734\nvalue = [0, 29]\nclass = b')]
```

```

lue = [5650, 201]\nclasse = a'),
  Text(0.6938775510204082, 0.5833333333333334, 'CO <= 0.55\ngini = 0.024\nsampl
les = 3469\nvalue = [5360, 67]\nclasse = a'),
  Text(0.6122448979591837, 0.4166666666666667, 'PM25 <= 28.5\ngini = 0.004\nsa
mples = 3109\nvalue = [4857, 10]\nclasse = a'),
  Text(0.5714285714285714, 0.25, 'NMHC <= 0.075\ngini = 0.003\nsamples = 3078
\nvalue = [4816, 7]\nclasse = a'),
  Text(0.5510204081632653, 0.0833333333333333, 'gini = 0.019\nsamples = 277\n
value = [424, 4]\nclasse = a'),
  Text(0.5918367346938775, 0.0833333333333333, 'gini = 0.001\nsamples = 2801
\nvalue = [4392, 3]\nclasse = a'),
  Text(0.6530612244897959, 0.25, 'TCH <= 1.445\ngini = 0.127\nsamples = 31\nva
lue = [41, 3]\nclasse = a'),
  Text(0.6326530612244898, 0.0833333333333333, 'gini = 0.337\nsamples = 10\nv
alue = [11, 3]\nclasse = a'),
  Text(0.673469387755102, 0.0833333333333333, 'gini = 0.0\nsamples = 21\nvalu
e = [30, 0]\nclasse = a'),
  Text(0.7755102040816326, 0.4166666666666667, 'SO_2 <= 5.5\ngini = 0.183\nsam
ples = 360\nvalue = [503, 57]\nclasse = a'),
  Text(0.7346938775510204, 0.25, 'EBE <= 0.65\ngini = 0.206\nsamples = 40\nval
ue = [7, 53]\nclasse = b'),
  Text(0.7142857142857143, 0.0833333333333333, 'gini = 0.083\nsamples = 30\nv
alue = [2, 44]\nclasse = b'),
  Text(0.7551020408163265, 0.0833333333333333, 'gini = 0.459\nsamples = 10\nv
alue = [5, 9]\nclasse = b'),
  Text(0.8163265306122449, 0.25, 'TCH <= 1.425\ngini = 0.016\nsamples = 320\nv
alue = [496, 4]\nclasse = a'),
  Text(0.7959183673469388, 0.0833333333333333, 'gini = 0.176\nsamples = 25\nv
alue = [37, 4]\nclasse = a'),
  Text(0.8367346938775511, 0.0833333333333333, 'gini = 0.0\nsamples = 295\nva
lue = [459, 0]\nclasse = a'),
  Text(0.9183673469387755, 0.5833333333333334, 'BEN <= 1.65\ngini = 0.432\nsam
ples = 265\nvalue = [290, 134]\nclasse = a'),
  Text(0.8775510204081632, 0.4166666666666667, 'PM25 <= 17.5\ngini = 0.246\nsa
mples = 90\nvalue = [22, 131]\nclasse = b'),
  Text(0.8571428571428571, 0.25, 'gini = 0.444\nsamples = 18\nvalue = [16, 8]
\nclasse = a'),
  Text(0.8979591836734694, 0.25, 'TCH <= 1.74\ngini = 0.089\nsamples = 72\nval
ue = [6, 123]\nclasse = b'),
  Text(0.8775510204081632, 0.0833333333333333, 'gini = 0.0\nsamples = 60\nval
ue = [0, 109]\nclasse = b'),
  Text(0.9183673469387755, 0.0833333333333333, 'gini = 0.42\nsamples = 12\nva
lue = [6, 14]\nclasse = b'),
  Text(0.9591836734693877, 0.4166666666666667, 'NO_2 <= 77.5\ngini = 0.022\nsa
mples = 175\nvalue = [268, 3]\nclasse = a'),
  Text(0.9387755102040817, 0.25, 'gini = 0.236\nsamples = 10\nvalue = [19, 3]
\nclasse = a'),
  Text(0.9795918367346939, 0.25, 'gini = 0.0\nsamples = 165\nvalue = [249, 0]
\nclasse = a')]

```



## Conclusion

```
In [34]: print("Linear Regression:", lr.score(x_test, y_test))
print("Ridge Regression:", rr.score(x_test, y_test))
print("Lasso Regression", la.score(x_test, y_test))
print("ElasticNet Regression:", en.score(x_test, y_test))
print("Logistic Regression:", logr.score(fs, target_vector))
print("Random Forest:", grid_search.best_score_)
```

Linear Regression: 0.8307069955786595  
 Ridge Regression: 0.8307958700138998  
 Lasso Regression 0.6552591684765237  
 ElasticNet Regression: 0.7176989302087358  
 Logistic Regression: 0.996161115048429  
 Random Forest: 0.9947688153898075

## Random Forest Is Better!!!