

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: df=pd.read_csv("madrid_2017.csv")
```

```
In [3]: df.head()
```

Out[3]:

	date	BEN	CH4	CO	EBE	NMHC	NO	NO_2	NOx	O_3	PM10	PM25	SO_2	TCH	TOL
0	2017-06-01 01:00:00	NaN	NaN	0.3	NaN	NaN	4.0	38.0	NaN	NaN	NaN	NaN	5.0	NaN	NaN
1	2017-06-01 01:00:00	0.6	NaN	0.3	0.4	0.08	3.0	39.0	NaN	71.0	22.0	9.0	7.0	1.4	NaN
2	2017-06-01 01:00:00	0.2	NaN	NaN	0.1	NaN	1.0	14.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN
3	2017-06-01 01:00:00	NaN	NaN	0.2	NaN	NaN	1.0	9.0	NaN	91.0	NaN	NaN	NaN	NaN	NaN
4	2017-06-01 01:00:00	NaN	NaN	NaN	NaN	NaN	1.0	19.0	NaN	69.0	NaN	NaN	2.0	NaN	NaN

```
In [4]: df=df.dropna()
```

```
In [5]: df.columns
```

Out[5]: Index(['date', 'BEN', 'CH4', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'NOx', 'O_3', 'PM10', 'PM25', 'SO_2', 'TCH', 'TOL', 'station'], dtype='object')

In [6]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 4127 entries, 87457 to 158286
Data columns (total 16 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   date        4127 non-null   object
 1   BEN         4127 non-null   float64
 2   CH4         4127 non-null   float64
 3   CO          4127 non-null   float64
 4   EBE         4127 non-null   float64
 5   NMHC        4127 non-null   float64
 6   NO          4127 non-null   float64
 7   NO_2        4127 non-null   float64
 8   NOx         4127 non-null   float64
 9   O_3         4127 non-null   float64
10  PM10        4127 non-null   float64
11  PM25        4127 non-null   float64
12  SO_2        4127 non-null   float64
13  TCH         4127 non-null   float64
14  TOL         4127 non-null   float64
15  station     4127 non-null   int64
dtypes: float64(14), int64(1), object(1)
memory usage: 548.1+ KB
```

In [7]: data=df[['CO', 'station']]
data

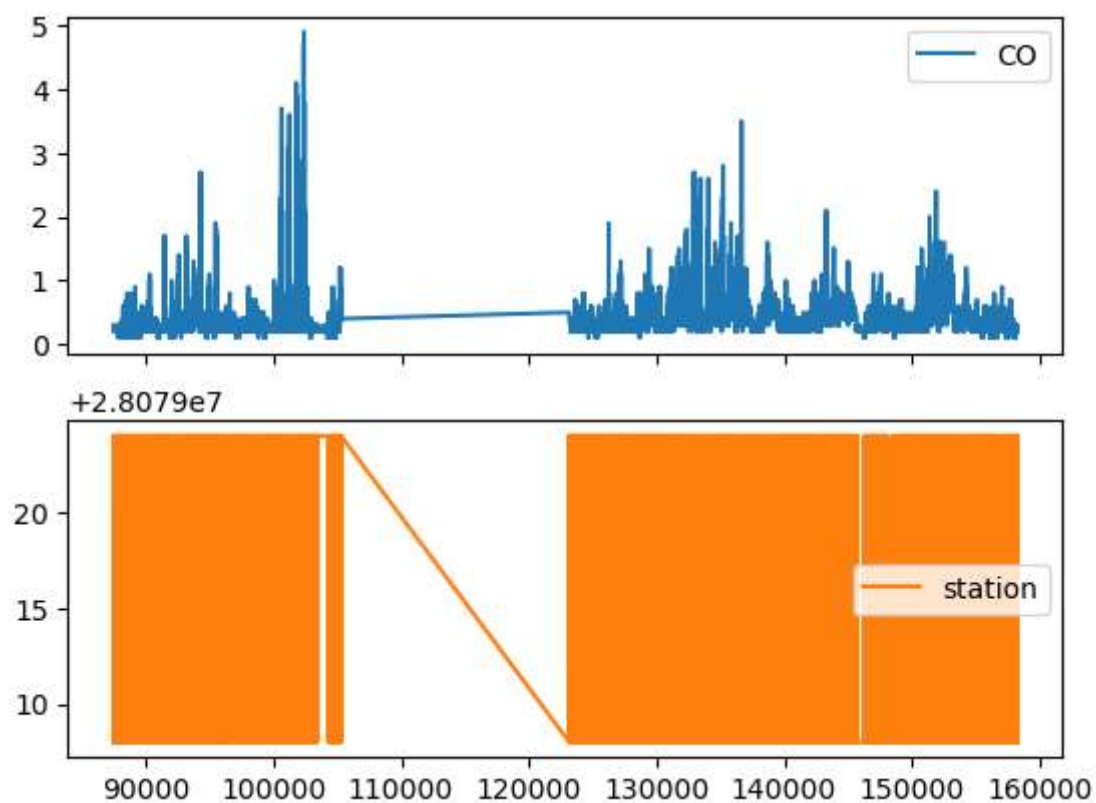
Out[7]:

	CO	station
87457	0.3	28079008
87462	0.2	28079024
87481	0.2	28079008
87486	0.2	28079024
87505	0.2	28079008
...
158238	0.2	28079024
158257	0.3	28079008
158262	0.2	28079024
158281	0.2	28079008
158286	0.2	28079024

4127 rows × 2 columns


```
In [8]: data.plot.line(subplots=True)
```

```
Out[8]: array([<Axes: >, <Axes: >], dtype=object)
```




```
In [9]: b=data[0:50]  
b.plot.bar()
```

Out[9]: <Axes: >




```
In [10]: data.plot.hist()
```

```
Out[10]: <Axes: ylabel='Frequency'>
```




```
In [11]: data.plot.area()
```

```
Out[11]: <Axes: >
```




```
In [12]: data.plot.box()
```

```
Out[12]: <Axes: >
```



```
In [13]: data.plot.scatter(x='CO',y='station')
```

```
Out[13]: <Axes: xlabel='CO', ylabel='station'>
```



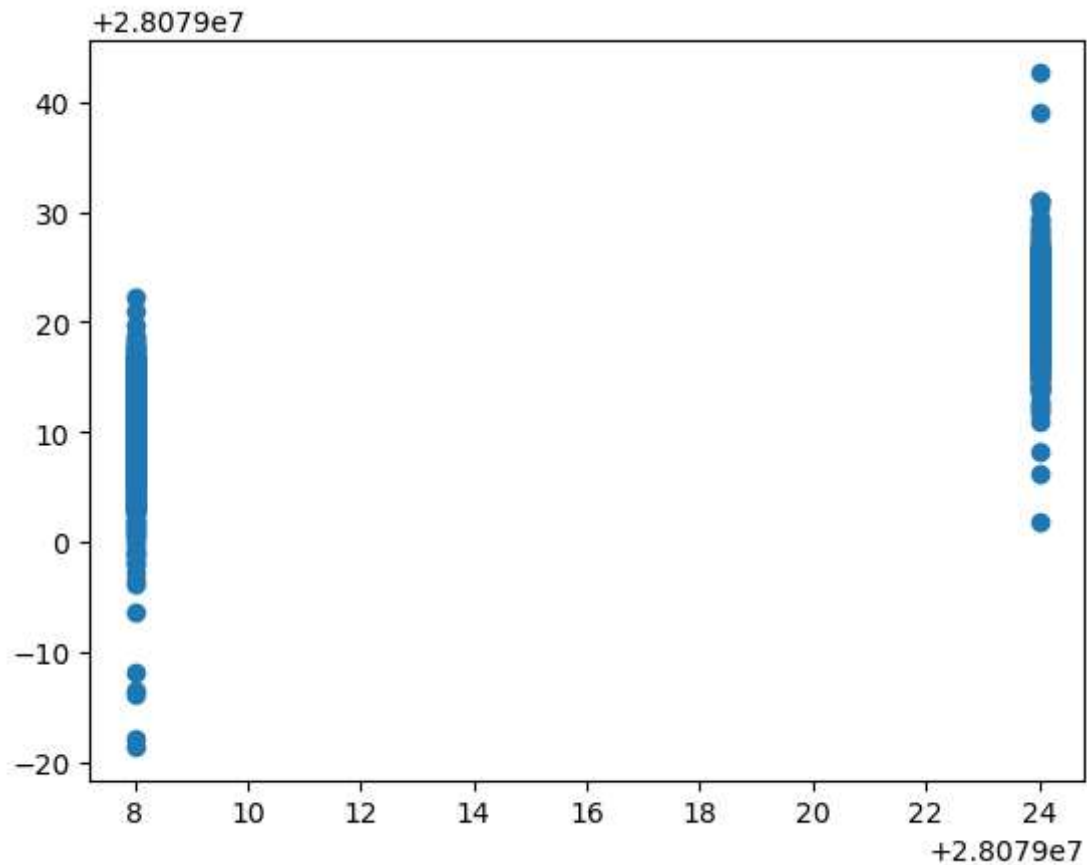
```
In [14]: x=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'NO', 'O_3',  
              'PM10','PM25','SO_2', 'TCH', 'TOL']]  
y=df['station']
```

```
In [15]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

Linear Regression


```
In [16]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
lr.intercept_
prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[16]: <matplotlib.collections.PathCollection at 0x19add634190>



```
In [17]: print(lr.score(x_test,y_test))
print(lr.score(x_train,y_train))
```

0.6266064332370529
0.6354532898583181

Ridge and Lasso

```
In [18]: from sklearn.linear_model import Ridge,Lasso
rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
print(rr.score(x_test,y_test))
print(rr.score(x_train,y_train))
la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

0.6267417967400755

0.6255645706561448

Out[18]:

▼ Lasso

Lasso(alpha=10)

```
In [19]: la.score(x_test,y_test)
```

Out[19]: 0.4007652853363365

ElasticNet

```
In [20]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

Out[20]:

▼ ElasticNet

ElasticNet()

```
In [21]: en.coef_
```

Out[21]: array([-0. , -0. , -0. , 0. , -0.21033405,
 0.03318752, -0.08640823, 0.52794605, -0.37262767, -0.31282917,
 -0. , 0.])

```
In [22]: en.intercept_
```

Out[22]: 28079025.708727796

```
In [23]: prediction=en.predict(x_test)
```

```
In [24]: en.score(x_test,y_test)
```

Out[24]: 0.5144988330589095

Evaluation Metrics

```
In [25]: from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
4.717827276519896
31.06865398929979
5.57392626335331
```

Logistics Regression

```
In [26]: from sklearn.linear_model import LogisticRegression
```

```
In [27]: feature_matrix=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'NO', 'O_3',
'PM10','PM25','SO_2', 'TCH', 'TOL']]
target_vector=df[ 'station']
```

```
In [28]: from sklearn.preprocessing import StandardScaler
fs=StandardScaler().fit_transform(feature_matrix)
logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)
```

```
Out[28]: LogisticRegression
LogisticRegression(max_iter=10000)
```

```
In [29]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12]]
logr.predict_proba(observation)
```

```
Out[29]: array([[1.00000000e+00, 1.00193473e-10]])
```

Random Forest

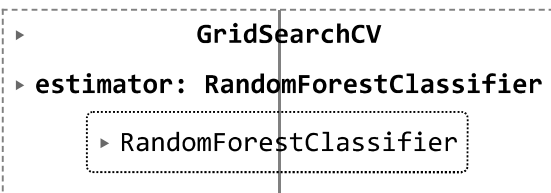
```
In [30]: from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[30]: RandomForestClassifier
RandomForestClassifier()
```

```
In [31]: parameters={'max_depth':[1,2,3,4,5],  
                    'min_samples_leaf':[5,10,15,20,25],  
                    'n_estimators':[10,20,30,40,50]  
                    }
```

```
In [32]: from sklearn.model_selection import GridSearchCV  
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="ac  
grid_search.fit(x_train,y_train)
```

```
Out[32]:
```



```
▶ GridSearchCV  
▶ estimator: RandomForestClassifier  
    ▶ RandomForestClassifier
```

```
In [33]: rfc_best=grid_search.best_estimator_  
         from sklearn.tree import plot_tree  
         plt.figure(figsize=(80,40))  
         plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b'])
```

```

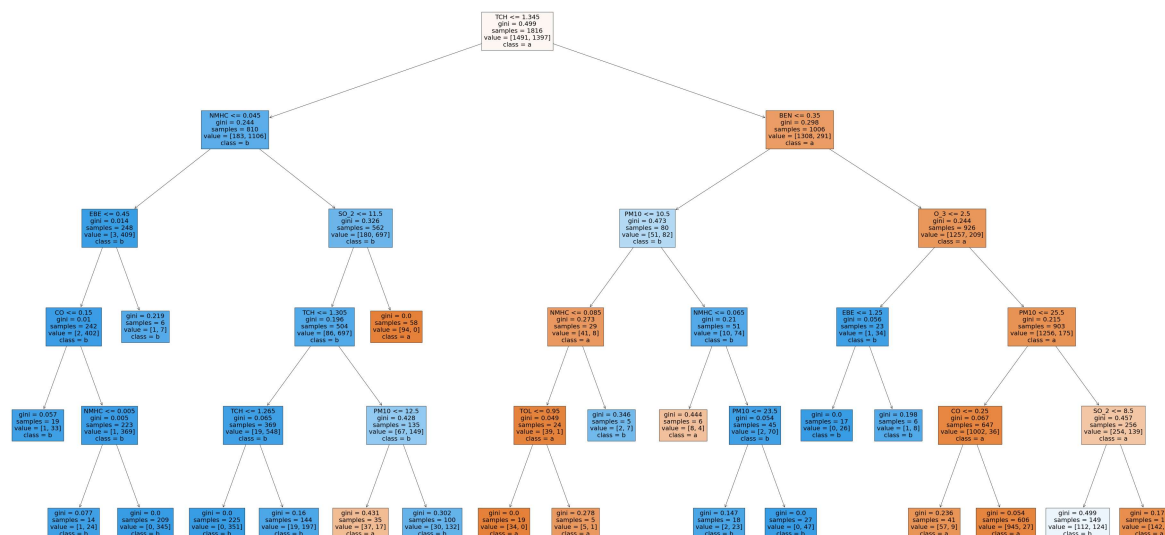
Out[33]: [Text(0.4356060606060606, 0.9166666666666666, 'TCH <= 1.345\ngini = 0.499\nsa
mples = 1816\nvalue = [1491, 1397]\nnclass = a'),
  Text(0.19696969696969696, 0.75, 'NMHC <= 0.045\ngini = 0.244\nsamples = 810
\nvalue = [183, 1106]\nnclass = b'),
  Text(0.09090909090909091, 0.5833333333333334, 'EBE <= 0.45\ngini = 0.014\nsa
mples = 248\nvalue = [3, 409]\nnclass = b'),
  Text(0.06060606060606061, 0.4166666666666667, 'CO <= 0.15\ngini = 0.01\nsamp
les = 242\nvalue = [2, 402]\nnclass = b'),
  Text(0.030303030303030304, 0.25, 'gini = 0.057\nsamples = 19\nvalue = [1, 3
3]\nnclass = b'),
  Text(0.09090909090909091, 0.25, 'NMHC <= 0.005\ngini = 0.005\nsamples = 223
\nvalue = [1, 369]\nnclass = b'),
  Text(0.06060606060606061, 0.08333333333333333, 'gini = 0.077\nsamples = 14\n
value = [1, 24]\nnclass = b'),
  Text(0.12121212121212122, 0.08333333333333333, 'gini = 0.0\nsamples = 209\nv
alue = [0, 345]\nnclass = b'),
  Text(0.12121212121212122, 0.4166666666666667, 'gini = 0.219\nsamples = 6\nva
lue = [1, 7]\nnclass = b'),
  Text(0.30303030303030304, 0.5833333333333334, 'SO_2 <= 11.5\ngini = 0.326\ns
amples = 562\nvalue = [180, 697]\nnclass = b'),
  Text(0.2727272727272727, 0.4166666666666667, 'TCH <= 1.305\ngini = 0.196\nsa
mples = 504\nvalue = [86, 697]\nnclass = b'),
  Text(0.21212121212121213, 0.25, 'TCH <= 1.265\ngini = 0.065\nsamples = 369\n
value = [19, 548]\nnclass = b'),
  Text(0.18181818181818182, 0.08333333333333333, 'gini = 0.0\nsamples = 225\nv
alue = [0, 351]\nnclass = b'),
  Text(0.24242424242424243, 0.08333333333333333, 'gini = 0.16\nsamples = 144\n
value = [19, 197]\nnclass = b'),
  Text(0.3333333333333333, 0.25, 'PM10 <= 12.5\ngini = 0.428\nsamples = 135\nv
alue = [67, 149]\nnclass = b'),
  Text(0.30303030303030304, 0.08333333333333333, 'gini = 0.431\nsamples = 35\n
value = [37, 17]\nnclass = a'),
  Text(0.36363636363636365, 0.08333333333333333, 'gini = 0.302\nsamples = 100
\nvalue = [30, 132]\nnclass = b'),
  Text(0.3333333333333333, 0.4166666666666667, 'gini = 0.0\nsamples = 58\nvalu
e = [94, 0]\nnclass = a'),
  Text(0.6742424242424242, 0.75, 'BEN <= 0.35\ngini = 0.298\nsamples = 1006\nv
alue = [1308, 291]\nnclass = a'),
  Text(0.5454545454545454, 0.5833333333333334, 'PM10 <= 10.5\ngini = 0.473\nsa
mples = 80\nvalue = [51, 82]\nnclass = b'),
  Text(0.48484848484848486, 0.4166666666666667, 'NMHC <= 0.085\ngini = 0.273\n
samples = 29\nvalue = [41, 8]\nnclass = a'),
  Text(0.45454545454545453, 0.25, 'TOL <= 0.95\ngini = 0.049\nsamples = 24\nva
lue = [39, 1]\nnclass = a'),
  Text(0.42424242424242425, 0.08333333333333333, 'gini = 0.0\nsamples = 19\nva
lue = [34, 0]\nnclass = a'),
  Text(0.48484848484848486, 0.08333333333333333, 'gini = 0.278\nsamples = 5\nv
alue = [5, 1]\nnclass = a'),
  Text(0.5151515151515151, 0.25, 'gini = 0.346\nsamples = 5\nvalue = [2, 7]\nc
lass = b'),
  Text(0.6060606060606061, 0.4166666666666667, 'NMHC <= 0.065\ngini = 0.21\nsa
mples = 51\nvalue = [10, 74]\nnclass = b'),
  Text(0.5757575757575758, 0.25, 'gini = 0.444\nsamples = 6\nvalue = [8, 4]\nc
lass = a'),
  Text(0.6363636363636364, 0.25, 'PM10 <= 23.5\ngini = 0.054\nsamples = 45\nva
lue = [2, 70]\nnclass = b'),
  Text(0.6060606060606061, 0.08333333333333333, 'gini = 0.147\nsamples = 18\nv

```

```

    value = [2, 23]\nnclass = b'),
    Text(0.6666666666666666, 0.08333333333333333, 'gini = 0.0\nsamples = 27\nvalue = [0, 47]\nnclass = b'),
    Text(0.803030303030303, 0.5833333333333334, 'O_3 <= 2.5\nngini = 0.244\nsamples = 926\nvalue = [1257, 209]\nnclass = a'),
    Text(0.7272727272727273, 0.4166666666666667, 'EBE <= 1.25\nngini = 0.056\nsamples = 23\nvalue = [1, 34]\nnclass = b'),
    Text(0.696969696969697, 0.25, 'gini = 0.0\nsamples = 17\nvalue = [0, 26]\nnclass = b'),
    Text(0.7575757575757576, 0.25, 'gini = 0.198\nsamples = 6\nvalue = [1, 8]\nnclass = b'),
    Text(0.8787878787878788, 0.4166666666666667, 'PM10 <= 25.5\nngini = 0.215\nsamples = 903\nvalue = [1256, 175]\nnclass = a'),
    Text(0.8181818181818182, 0.25, 'CO <= 0.25\nngini = 0.067\nsamples = 647\nvalue = [1002, 36]\nnclass = a'),
    Text(0.7878787878787878, 0.08333333333333333, 'gini = 0.236\nsamples = 41\nvalue = [57, 9]\nnclass = a'),
    Text(0.8484848484848485, 0.08333333333333333, 'gini = 0.054\nsamples = 606\nvalue = [945, 27]\nnclass = a'),
    Text(0.9393939393939394, 0.25, 'SO_2 <= 8.5\nngini = 0.457\nsamples = 256\nvalue = [254, 139]\nnclass = a'),
    Text(0.9090909090909091, 0.08333333333333333, 'gini = 0.499\nsamples = 149\nvalue = [112, 124]\nnclass = b'),
    Text(0.9696969696969697, 0.08333333333333333, 'gini = 0.173\nsamples = 107\nvalue = [142, 15]\nnclass = a'))

```



Conclusion

```
In [34]: print("Linear Regression:",lr.score(x_test,y_test))
print("Ridge Regression:",rr.score(x_test,y_test))
print("Lasso Regression",la.score(x_test,y_test))
print("ElasticNet Regression:",en.score(x_test,y_test))
print("Logistic Regression:",logr.score(fs,target_vector))
print("Random Forest:",grid_search.best_score_)
```

Linear Regression: 0.6266064332370529
Ridge Regression: 0.6267417967400755
Lasso Regression 0.4007652853363365
ElasticNet Regression: 0.5144988330589095
Logistic Regression: 0.9520232614489944
Random Forest: 0.96398891966759

Random Forest Is Better!!!