

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: df=pd.read_csv("madrid_2018.csv")
```

```
In [3]: df.head()
```

Out[3]:

	date	BEN	CH4	CO	EBE	NMHC	NO	NO_2	NOx	O_3	PM10	PM25	SO_2	TCH	TOL	station
0	2018-03-01 01:00:00	NaN	NaN	0.3	NaN	NaN	1.0	29.0	31.0	NaN	NaN	NaN	2.0	NaN	NaN	1
1	2018-03-01 01:00:00	0.5	1.39	0.3	0.2	0.02	6.0	40.0	49.0	52.0	5.0	4.0	3.0	1.41	NaN	1
2	2018-03-01 01:00:00	0.4	NaN	NaN	0.2	NaN	4.0	41.0	47.0	NaN	NaN	NaN	NaN	NaN	NaN	1
3	2018-03-01 01:00:00	NaN	NaN	0.3	NaN	NaN	1.0	35.0	37.0	54.0	NaN	NaN	NaN	NaN	NaN	1
4	2018-03-01 01:00:00	NaN	NaN	NaN	NaN	NaN	1.0	27.0	29.0	49.0	NaN	NaN	3.0	NaN	NaN	1

```
In [4]: df=df.dropna()
```

```
In [5]: df.columns
```

Out[5]: Index(['date', 'BEN', 'CH4', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'NOx', 'O_3', 'PM10', 'PM25', 'SO_2', 'TCH', 'TOL', 'station'], dtype='object')

In [6]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 4562 entries, 1 to 69078
Data columns (total 16 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        4562 non-null   object
1   BEN         4562 non-null   float64
2   CH4         4562 non-null   float64
3   CO          4562 non-null   float64
4   EBE         4562 non-null   float64
5   NMHC        4562 non-null   float64
6   NO          4562 non-null   float64
7   NO_2        4562 non-null   float64
8   NOx         4562 non-null   float64
9   O_3         4562 non-null   float64
10  PM10        4562 non-null   float64
11  PM25        4562 non-null   float64
12  SO_2        4562 non-null   float64
13  TCH         4562 non-null   float64
14  TOL         4562 non-null   float64
15  station     4562 non-null   int64
dtypes: float64(14), int64(1), object(1)
memory usage: 605.9+ KB
```

In [7]: data=df[['CO', 'station']]
data


Out[7]:

	CO	station
1	0.3	28079008
6	0.2	28079024
25	0.2	28079008
30	0.2	28079024
49	0.2	28079008
...
69030	0.7	28079024
69049	1.2	28079008
69054	0.6	28079024
69073	1.0	28079008
69078	0.4	28079024

4562 rows × 2 columns


```
In [8]: data.plot.line(subplots=True)
```

```
Out[8]: array([<Axes: >, <Axes: >], dtype=object)
```




```
In [9]: b=data[0:50]  
b.plot.bar()
```

Out[9]: <Axes: >




```
In [10]: data.plot.hist()
```

```
Out[10]: <Axes: ylabel='Frequency'>
```




```
In [11]: data.plot.area()
```

```
Out[11]: <Axes: >
```




```
In [12]: data.plot.box()
```

```
Out[12]: <Axes: >
```



```
In [13]: data.plot.scatter(x='CO',y='station')
```

```
Out[13]: <Axes: xlabel='CO', ylabel='station'>
```




```
In [14]: x=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'NO', 'O_3',  
              'PM10','PM25','SO_2', 'TCH', 'TOL']]  
y=df['station']
```

```
In [15]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

Linear Regression


```
In [16]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
lr.intercept_
prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[16]: <matplotlib.collections.PathCollection at 0x1898cbfab90>



```
In [17]: print(lr.score(x_test,y_test))
print(lr.score(x_train,y_train))
```

```
0.8244767636678925
0.8027984588266978
```

Ridge and Lasso

```
In [18]: from sklearn.linear_model import Ridge,Lasso
rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
print(rr.score(x_test,y_test))
print(rr.score(x_train,y_train))
la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

0.7150863943178736

0.6992780061799038

Out[18]:

▼ Lasso

Lasso(alpha=10)

```
In [19]: la.score(x_test,y_test)
```

Out[19]: 0.41439299623014336

ElasticNet

```
In [20]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

Out[20]:

▼ ElasticNet

ElasticNet()

```
In [21]: en.coef_
```

Out[21]: array([-0. , -0. , -0. , 0. , -0.28845039,
 0.03420383, -0.14450119, 0.27151285, -0.06886965, 0.06094714,
 -0.10779873, 0.])

```
In [22]: en.intercept_
```

Out[22]: 28079029.799183168

```
In [23]: prediction=en.predict(x_test)
```

```
In [24]: en.score(x_test,y_test)
```

Out[24]: 0.4597247930319087

Evaluation Metrics

```
In [25]: from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
4.918455148925304
33.18235851505395
5.760413050732903
```

Logistics Regression

```
In [26]: from sklearn.linear_model import LogisticRegression
```

```
In [27]: feature_matrix=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'NO', 'O_3',
'PM10','PM25','SO_2', 'TCH', 'TOL']]
target_vector=df[ 'station']
```

```
In [28]: from sklearn.preprocessing import StandardScaler
fs=StandardScaler().fit_transform(feature_matrix)
logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)
```

```
Out[28]: LogisticRegression
LogisticRegression(max_iter=10000)
```

```
In [29]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12]]
logr.predict_proba(observation)
```

```
Out[29]: array([[1.00000000e+00, 1.57204801e-20]])
```

Random Forest

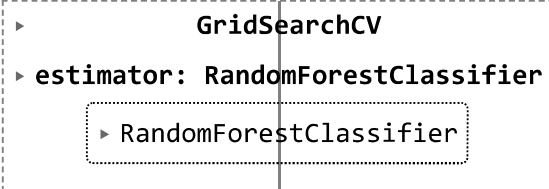
```
In [30]: from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[30]: RandomForestClassifier
RandomForestClassifier()
```

```
In [31]: parameters={'max_depth':[1,2,3,4,5],  
                  'min_samples_leaf':[5,10,15,20,25],  
                  'n_estimators':[10,20,30,40,50]  
                }
```

```
In [32]: from sklearn.model_selection import GridSearchCV  
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="ac  
grid_search.fit(x_train,y_train)
```

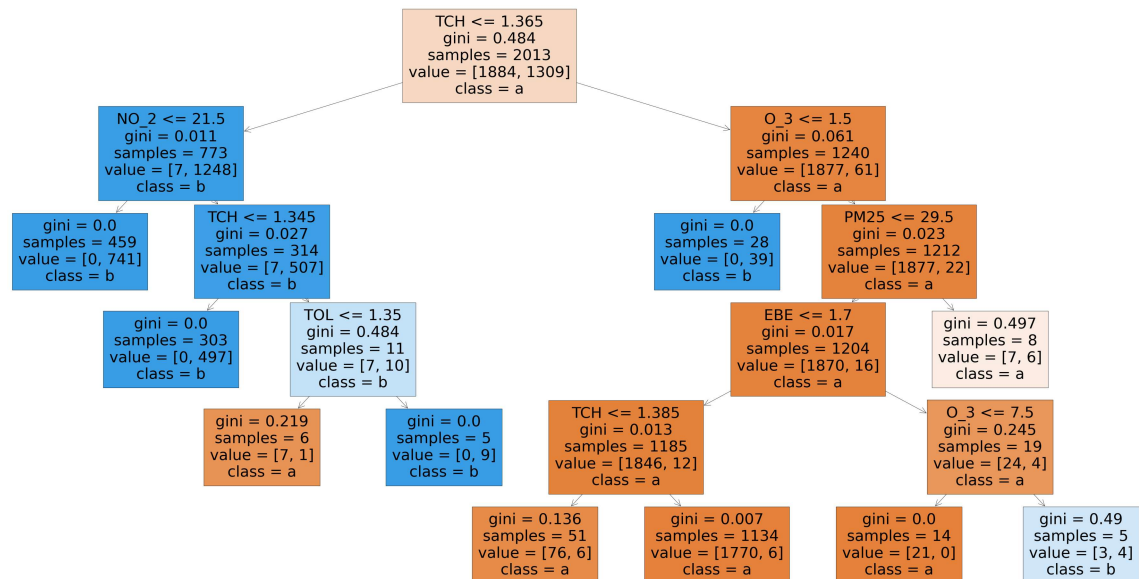
```
Out[32]:
```



```
▶ GridSearchCV  
▶ estimator: RandomForestClassifier  
    ▶ RandomForestClassifier
```

```
In [33]: rfc_best=grid_search.best_estimator_
from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b'])
```

```
Out[33]: [Text(0.4230769230769231, 0.9166666666666666, 'TCH <= 1.365\ngini = 0.484\nsamples = 2013\nvalue = [1884, 1309]\nnclass = a'),
Text(0.15384615384615385, 0.75, 'NO_2 <= 21.5\ngini = 0.011\nsamples = 773\nvalue = [7, 1248]\nnclass = b'),
Text(0.07692307692307693, 0.5833333333333334, 'gini = 0.0\nsamples = 459\nvalue = [0, 741]\nnclass = b'),
Text(0.23076923076923078, 0.5833333333333334, 'TCH <= 1.345\ngini = 0.027\nsamples = 314\nvalue = [7, 507]\nnclass = b'),
Text(0.15384615384615385, 0.4166666666666667, 'gini = 0.0\nsamples = 303\nvalue = [0, 497]\nnclass = b'),
Text(0.3076923076923077, 0.4166666666666667, 'TOL <= 1.35\ngini = 0.484\nsamples = 11\nvalue = [7, 10]\nnclass = b'),
Text(0.23076923076923078, 0.25, 'gini = 0.219\nsamples = 6\nvalue = [7, 1]\nnclass = a'),
Text(0.38461538461538464, 0.25, 'gini = 0.0\nsamples = 5\nvalue = [0, 9]\nnclass = b'),
Text(0.6923076923076923, 0.75, 'O_3 <= 1.5\ngini = 0.061\nsamples = 1240\nvalue = [1877, 61]\nnclass = a'),
Text(0.6153846153846154, 0.5833333333333334, 'gini = 0.0\nsamples = 28\nvalue = [0, 39]\nnclass = b'),
Text(0.7692307692307693, 0.5833333333333334, 'PM25 <= 29.5\ngini = 0.023\nsamples = 1212\nvalue = [1877, 22]\nnclass = a'),
Text(0.6923076923076923, 0.4166666666666667, 'EBE <= 1.7\ngini = 0.017\nsamples = 1204\nvalue = [1870, 16]\nnclass = a'),
Text(0.5384615384615384, 0.25, 'TCH <= 1.385\ngini = 0.013\nsamples = 1185\nvalue = [1846, 12]\nnclass = a'),
Text(0.46153846153846156, 0.08333333333333333, 'gini = 0.136\nsamples = 51\nvalue = [76, 6]\nnclass = a'),
Text(0.6153846153846154, 0.08333333333333333, 'gini = 0.007\nsamples = 1134\nvalue = [1770, 6]\nnclass = a'),
Text(0.8461538461538461, 0.25, 'O_3 <= 7.5\ngini = 0.245\nsamples = 19\nvalue = [24, 4]\nnclass = a'),
Text(0.7692307692307693, 0.08333333333333333, 'gini = 0.0\nsamples = 14\nvalue = [21, 0]\nnclass = a'),
Text(0.9230769230769231, 0.08333333333333333, 'gini = 0.49\nsamples = 5\nvalue = [3, 4]\nnclass = b'),
Text(0.8461538461538461, 0.4166666666666667, 'gini = 0.497\nsamples = 8\nvalue = [7, 6]\nnclass = a')]
```



Conclusion

```
In [34]: print("Linear Regression:",lr.score(x_test,y_test))
print("Ridge Regression:",rr.score(x_test,y_test))
print("Lasso Regression",la.score(x_test,y_test))
print("ElasticNet Regression:",en.score(x_test,y_test))
print("Logistic Regression:",logr.score(fs,target_vector))
print("Random Forest:",grid_search.best_score_)
```

Linear Regression: 0.8244767636678925
 Ridge Regression: 0.7150863943178736
 Lasso Regression 0.41439299623014336
 ElasticNet Regression: 0.4597247930319087
 Logistic Regression: 0.9890398947829899
 Random Forest: 0.9934232104996368

Random Forest Is Better!!!